

D212_Task_2

April 8, 2024

A1: PCA Question Which principal factors contribute the most variance in our customer data, and can these factors provide insights into characteristics associated with customer churn?

A2: Goal of the Data Analysis The goal of this analysis is to reduce the dimensionality of our customer dataset, thereby identifying key variables that account for the majority of variance. This will help in uncovering patterns that could be associated with churn, providing a basis for more targeted customer retention strategies to the executives.

B1: PCA Analysis Explanation PCA will work by converting my original set of possibly correlated variables into a smaller set of uncorrelated variables called principal components. The primary goal of this analysis is to utilize Principal Component Analysis (PCA) for dimensionality reduction within our churn dataset. By identifying the principal factors that account for the majority of variance, we aim to simplify the dataset, making it more manageable for further analysis. Recognizing PCA's limitations, I understand that while it does not directly predict or reveal specific patterns associated with customer churn, it effectively identifies key variables that merit further investigation through targeted analytical models. Subsequently, insights derived from these principal factors will guide the development of more nuanced customer retention strategies by highlighting potential areas of focus. This step is foundational, setting the stage for applying predictive modeling techniques to directly examine the relationship between these key factors and customer churn.

Jolliffe, I. T., & Cadima, J. (2016).

B2: Assumption of PCA One key assumption of PCA is that linear relationships exist among the variables. PCA assumes that the principal components are a linear combination of the original features and that these components capture the underlying structure of the data through linear correlations. This assumption is critical for the effective application of PCA, as it relies on maximizing variance through linear transformations.

Jolliffe, I. T., & Cadima, J. (2016).

```
[19]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA

# Loading the dataset
```

```
df = pd.read_csv(r'C:\Users\Hien Ta\OneDrive\WGU\MSDA\D212\Task_2\churn_clean.
↳csv')
```

```
# Check data types, number of values and size of dataframe
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 10000 entries, 0 to 9999
```

```
Data columns (total 50 columns):
```

#	Column	Non-Null Count	Dtype
0	CaseOrder	10000 non-null	int64
1	Customer_id	10000 non-null	object
2	Interaction	10000 non-null	object
3	UID	10000 non-null	object
4	City	10000 non-null	object
5	State	10000 non-null	object
6	County	10000 non-null	object
7	Zip	10000 non-null	int64
8	Lat	10000 non-null	float64
9	Lng	10000 non-null	float64
10	Population	10000 non-null	int64
11	Area	10000 non-null	object
12	TimeZone	10000 non-null	object
13	Job	10000 non-null	object
14	Children	10000 non-null	int64
15	Age	10000 non-null	int64
16	Income	10000 non-null	float64
17	Marital	10000 non-null	object
18	Gender	10000 non-null	object
19	Churn	10000 non-null	object
20	Outage_sec_perweek	10000 non-null	float64
21	Email	10000 non-null	int64
22	Contacts	10000 non-null	int64
23	Yearly_equip_failure	10000 non-null	int64
24	Techie	10000 non-null	object
25	Contract	10000 non-null	object
26	Port_modem	10000 non-null	object
27	Tablet	10000 non-null	object
28	InternetService	7871 non-null	object
29	Phone	10000 non-null	object
30	Multiple	10000 non-null	object
31	OnlineSecurity	10000 non-null	object
32	OnlineBackup	10000 non-null	object
33	DeviceProtection	10000 non-null	object
34	TechSupport	10000 non-null	object
35	StreamingTV	10000 non-null	object
36	StreamingMovies	10000 non-null	object

```

37 PaperlessBilling      10000 non-null object
38 PaymentMethod        10000 non-null object
39 Tenure                10000 non-null float64
40 MonthlyCharge         10000 non-null float64
41 Bandwidth_GB_Year    10000 non-null float64
42 Item1                10000 non-null int64
43 Item2                10000 non-null int64
44 Item3                10000 non-null int64
45 Item4                10000 non-null int64
46 Item5                10000 non-null int64
47 Item6                10000 non-null int64
48 Item7                10000 non-null int64
49 Item8                10000 non-null int64
dtypes: float64(7), int64(16), object(27)
memory usage: 3.8+ MB

```

```

[20]: # Identifying Missing Values
missing_values = df.isnull().sum()

print(missing_values)

```

```

CaseOrder      0
Customer_id    0
Interaction     0
UID            0
City           0
State          0
County         0
Zip            0
Lat            0
Lng            0
Population     0
Area           0
TimeZone       0
Job            0
Children       0
Age            0
Income         0
Marital        0
Gender         0
Churn          0
Outage_sec_perweek 0
Email          0
Contacts       0
Yearly_equip_failure 0
Techie         0
Contract       0
Port_modem     0

```

Tablet	0
InternetService	2129
Phone	0
Multiple	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
PaperlessBilling	0
PaymentMethod	0
Tenure	0
MonthlyCharge	0
Bandwidth_GB_Year	0
Item1	0
Item2	0
Item3	0
Item4	0
Item5	0
Item6	0
Item7	0
Item8	0

dtype: int64

```
[21]: # Fill missing value in 'InternetService' with 'None'
df['InternetService'].fillna('None', inplace=True)
```

```
[22]: # Identifying Missing Values
missing_values = df.isnull().sum()

print(missing_values)
```

CaseOrder	0
Customer_id	0
Interaction	0
UID	0
City	0
State	0
County	0
Zip	0
Lat	0
Lng	0
Population	0
Area	0
TimeZone	0
Job	0
Children	0

Age	0
Income	0
Marital	0
Gender	0
Churn	0
Outage_sec_perweek	0
Email	0
Contacts	0
Yearly_equip_failure	0
Techie	0
Contract	0
Port_modem	0
Tablet	0
InternetService	0
Phone	0
Multiple	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
PaperlessBilling	0
PaymentMethod	0
Tenure	0
MonthlyCharge	0
Bandwidth_GB_Year	0
Item1	0
Item2	0
Item3	0
Item4	0
Item5	0
Item6	0
Item7	0
Item8	0

dtype: int64

C1: Continuous Variables For the PCA analysis to address the question of identifying customer segments at the highest risk of churn and their common characteristics, I selected the following continuous variables from the dataset:

- Population: Reflects the population within the dataset.
- Age: The age of the customer.
- Income: The annual income of the customer.
- Outage_sec_perweek: Average seconds per week of system outages in the customer's area.
- Email: The number of emails sent to the customer in the last year, including both marketing and correspondence.
- Contacts: How many times the customer contacted technical support.

- Yearly_equip_failure: The number of times the customer's equipment failed and needed resetting/replacement in the past year.
- Tenure: The number of months the customer has been with the provider.
- MonthlyCharge: The average monthly amount charged to the customer. e* Bandwidth_GB_Year: The average annual data usage in GB by the customer. For new customers.

C2: Standardize Variable and Saving Dataset

```
[23]: # Selecting the continuous variables for PCA
continuous_vars = ['Population', 'Age', 'Income', 'Outage_sec_perweek',
↳ 'Email', 'Contacts',
                    'Yearly_equip_failure', 'Tenure', 'MonthlyCharge',
↳ 'Bandwidth_GB_Year']

# Standardizing the continuous variables
scaler = StandardScaler()
data_standardized = scaler.fit_transform(df[continuous_vars])

# Converting the standardized data back to a dataframe for readability
data_standardized_df = pd.DataFrame(data_standardized, columns=continuous_vars)

# Adding the feature names
feature_names = df[continuous_vars].columns.tolist()

# Saving the cleaned dataset
cleaned_file_path = (r'C:\Users\Hien_
↳ Ta\OneDrive\WGU\MSDA\D212\Task_2\churn_clean_After.csv')
data_standardized_df.to_csv(cleaned_file_path, index=False)
```

D1: Matrix Of All The Principal Components

```
[26]: # Initialize PCA
pca = PCA()

# Fit PCA on the standardized data
pca.fit(data_standardized)

# The matrix of all principal components
principal_components_matrix = pca.components_

# Kaiser criterion (eigenvalues > 1)
num_components_kaiser = np.sum(pca.explained_variance_ > 1)

print(f"Number of principal components to keep based on the Kaiser criterion:
↳ {num_components_kaiser}")
print('\n')
```

```

# Convert the matrix of principal components to a DataFrame
principal_components_df = pd.DataFrame( principal_components_matrix,
                                         columns=feature_names, # Use the
                                         ↪feature names for the columns
                                         index=[f'PC_{i+1}' for i in
                                         ↪range(principal_components_matrix.shape[0])])

# Transform the standardized data to principal components
data_in_principal_components = pca.transform(data_standardized)

# Convert the transformed data to a DataFrame with principal components as
↪columns
transformed_data_df = pd.DataFrame(data_in_principal_components,
                                   columns=[f'PC_{i+1}' for i in
                                   ↪range(data_in_principal_components.shape[1])])

print(principal_components_df.round(2))
print('\n')
print("-" * 100)
print('\n')
print(transformed_data_df.head().round(2)) # Display the first five rows of
↪the transformed data

# Scikit-learn developers. (n.d.)
# W3Schools.com. (n.d.)

```

Number of principal components to keep based on the Kaiser criterion: 5

	Population	Age	Income	Outage_sec_perweek	Email	Contacts	\
PC_1	-0.01	0.00	0.00	0.01	-0.02	0.00	
PC_2	-0.39	-0.22	0.34	-0.40	-0.44	-0.36	
PC_3	0.37	-0.18	-0.15	-0.42	0.44	-0.32	
PC_4	0.32	0.72	0.08	-0.34	-0.06	0.26	
PC_5	-0.23	0.09	0.57	-0.32	0.10	0.40	
PC_6	0.25	-0.40	0.31	0.38	0.01	0.48	
PC_7	0.37	-0.01	0.63	0.12	0.29	-0.41	
PC_8	0.56	-0.05	0.03	-0.00	-0.72	-0.10	
PC_9	0.20	-0.48	-0.16	-0.54	0.02	0.37	
PC_10	-0.00	0.02	-0.00	0.00	0.00	-0.00	

	Yearly_equip_failure	Tenure	MonthlyCharge	Bandwidth_GB_Year
PC_1	0.02	0.71	0.04	0.71
PC_2	0.29	0.01	-0.35	-0.01
PC_3	-0.33	0.04	-0.46	0.02
PC_4	0.36	0.02	-0.24	-0.02
PC_5	-0.57	-0.00	0.12	0.01

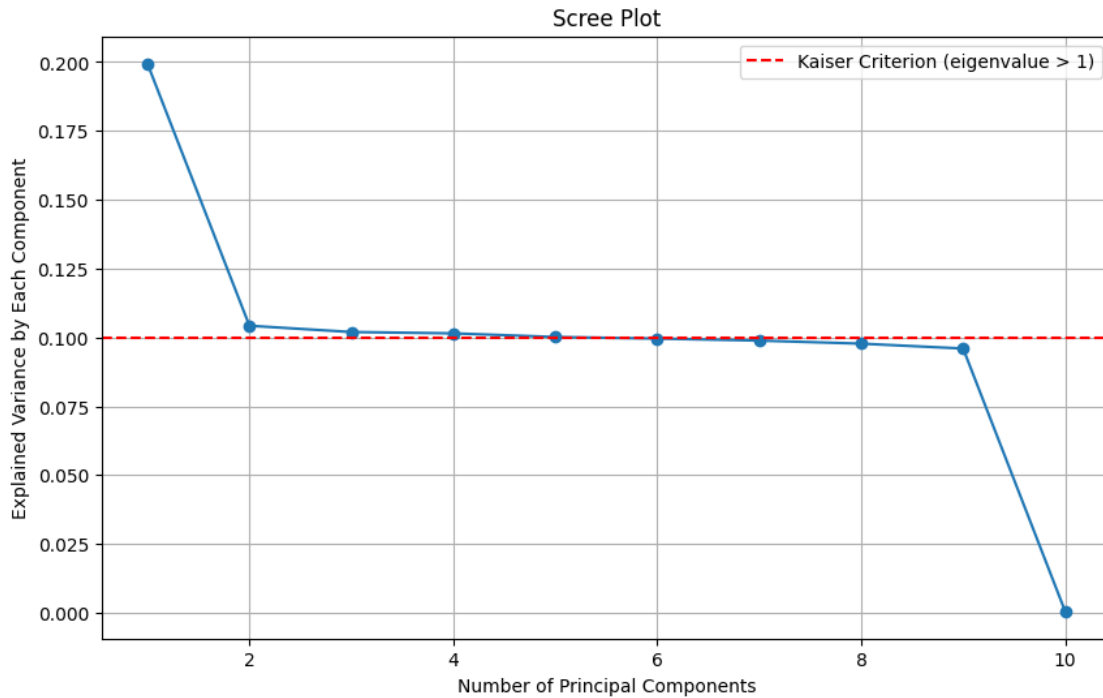
PC_6	0.16	0.02	-0.53	0.00
PC_7	0.26	-0.02	0.35	0.00
PC_8	-0.38	-0.01	0.11	-0.00
PC_9	0.33	-0.03	0.41	0.01
PC_10	-0.00	-0.71	-0.05	0.71

	PC_1	PC_2	PC_3	PC_4	PC_5	PC_6	PC_7	PC_8	PC_9	PC_10
0	-1.52	1.17	-0.38	0.62	-0.81	-1.19	-0.11	-0.20	-0.10	-0.05
1	-1.65	-0.13	-0.73	-1.45	-1.42	-0.69	0.95	0.00	1.06	-0.05
2	-0.91	0.90	-0.41	-0.21	-1.69	-0.48	-0.52	0.17	-0.17	0.08
3	-0.94	-1.53	0.42	-0.42	-0.32	1.59	-0.88	-0.54	-0.89	0.12
4	-1.92	-0.58	0.16	1.95	0.22	0.12	-0.01	-1.44	0.19	-0.08

D2: Total Number of Principal Components

```
[28]: # Generate scree plot
plt.figure(figsize=(10, 6))
plt.plot(range(1, len(pca.explained_variance_ratio_) + 1), pca.
         explained_variance_ratio_, marker='o')
plt.title('Scree Plot')
plt.xlabel('Number of Principal Components')
plt.ylabel('Explained Variance by Each Component')
plt.axhline(y=1/len(pca.explained_variance_ratio_), color='r', linestyle='--',
            label='Kaiser Criterion (eigenvalue > 1)')
plt.grid(True)
plt.legend()
plt.show()

# Scikit-learn developers. (n.d.)
# W3Schools.com. (n.d.)
```

D3: Variance of each of the principal components

```
[29]: # Extract the variance explained by the number of principal components as
      ↪determined by the Kaiser criterion
      variance_by_component = pca.explained_variance_ratio_[:num_components_kaiser]
      for i, variance in enumerate(variance_by_component, 1):
          print(f"PC{i}: {variance:.4f} or {variance * 100:.2f}% of the total
          ↪variance")
```

PC1: 0.1994 or 19.94% of the total variance

PC2: 0.1042 or 10.42% of the total variance

PC3: 0.1020 or 10.20% of the total variance

PC4: 0.1015 or 10.15% of the total variance

PC5: 0.1002 or 10.02% of the total variance

D4: Identify the Total Variance Captured by the Principal Components Identified in Part D2

```
[30]: # Calculate the total variance captured by the first 5 principal components
      total_variance_captured = np.sum(variance_by_component)
      print(f"\nTotal variance captured by the first 5 components:
      ↪{total_variance_captured:.4f} or {total_variance_captured * 100:.2f}%")
```

Total variance captured by the first 5 components: 0.6072 or 60.72%

D5: Summary of Analysis Based on the Kaiser criterion, I retained the first 5 principal components for my PCA analysis. These components together explain 60.72% of the total variance in the dataset. This indicate a significant reduction in dimensionality while still capturing a substantial amount of information. The variance explained by each of the components highlights their relative importance, with PC1 explaining the most variance and PC5 the least of the selected components.

This reduction enables me to focus on the most meaningful patterns in the data, which can be particularly useful for further analysis. For example clustering or predictive modeling. By understanding the major drivers of variance captured by these components, we can gain insights into underlying factors that might be influencing customer behavior. I have pinpointed the most influential variables shaping customer behavior. This reduction highlights critical factors like tenure, monthlycharges, and usage patterns, suggesting their significant role in customer differentiation and potential churn. By analyzing these components, I identify key customer attributes that contribute to variance, enabling targeted retention strategies. Furthermore by providing a foundation for deeper insights into customer behavior and guiding more effective decision making to mitigate churn risk. This analysis not only aids in simplifying complex datasets, but also in uncovering potential areas for strategic decision making for the executives.

E1: Third-Party Source No third-party source used

F1: Source Jolliffe, I. T., & Cadima, J. (2016). Principal component analysis: a review and recent developments. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 374(2065), 20150202. <https://doi.org/10.1098/rsta.2015.0202>

Scikit-learn developers. (n.d.). Principal component analysis (PCA). scikit-learn. Retrieved April 3, 2024, from <https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.PCA.html>

W3Schools.com. (n.d.). Python Machine Learning - Getting Started. Retrieved April 3, 2024, from https://www.w3schools.com/python/python_ml_getting_started.asp